

Assignment 3: Medical Image Segmentation with Deep Learning

Neural Networks and Deep Learning
U-Net and Attention U-Net for Brain Tissue Segmentation

Student Name

February 13, 2026

Abstract

This report presents a comprehensive implementation and evaluation of deep learning models for medical image segmentation, specifically focusing on brain tissue segmentation using the IBSR (Internet Brain Segmentation Repository) dataset. We implement and compare two state-of-the-art architectures: the classical U-Net and an enhanced Attention U-Net. The U-Net architecture demonstrates excellent performance in medical image segmentation tasks due to its encoder-decoder structure with skip connections, which preserves fine-grained spatial information. The Attention U-Net extends this by incorporating attention gates in the skip connections, allowing the model to focus on relevant features and suppress irrelevant background information. Our implementation includes comprehensive data preprocessing, model training with appropriate loss functions (combined Dice and Cross-Entropy loss), and extensive evaluation using multiple metrics including Dice coefficient, Intersection over Union (IoU), and pixel accuracy. The results demonstrate the effectiveness of both architectures, with detailed analysis of model performance, error patterns, and visualization of predictions. This work provides a complete pipeline for medical image segmentation that can be adapted for various clinical applications.

Contents

1 Introduction

2 Introduction

Medical image segmentation represents one of the most critical and challenging tasks in computer-aided diagnosis, surgical planning, and treatment monitoring. The ability to accurately delineate anatomical structures from medical images enables precise quantification of tissue volumes, disease progression tracking, treatment response assessment, and surgical navigation. In clinical practice, segmentation accuracy directly impacts patient outcomes, making it a cornerstone of modern medical imaging workflows.

2.1 The Critical Role of Medical Image Segmentation

Medical image segmentation serves as the foundation for numerous clinical applications across multiple medical domains:

2.1.1 Clinical Applications and Impact

- **Cancer Diagnosis and Staging:** Accurate tumor segmentation enables precise volume measurements, staging determination, and treatment planning. For instance, in brain tumor segmentation, sub-millimeter accuracy can determine surgical resectability and influence treatment protocols.
- **Cardiovascular Assessment:** Segmentation of cardiac chambers and vessels from MRI/CT images enables ejection fraction calculation, wall motion analysis, and coronary artery disease assessment.
- **Neurological Disorders:** Brain tissue segmentation facilitates the quantification of atrophy in neurodegenerative diseases like Alzheimer's, multiple sclerosis lesion tracking, and stroke volume assessment.
- **Orthopedic Applications:** Bone segmentation from CT scans enables implant design, fracture assessment, and joint space analysis for arthritis monitoring.
- **Radiotherapy Planning:** Precise organ and tumor segmentation ensures accurate radiation dose delivery while sparing healthy tissues.

2.1.2 Economic and Healthcare System Impact

Beyond individual patient care, accurate segmentation contributes significantly to healthcare economics by:

- Reducing diagnostic interpretation time by up to 70% through automated preprocessing
- Improving treatment planning accuracy, potentially reducing costly re-treatments
- Enabling population-level disease monitoring and epidemiological studies
- Supporting telemedicine applications in underserved regions

2.2 Unique Challenges in Medical Image Segmentation

Medical image segmentation presents distinct challenges that differentiate it from natural image segmentation tasks:

2.2.1 Imaging Modality-Specific Challenges

Magnetic Resonance Imaging (MRI):

- **Intensity Inhomogeneity:** RF coil sensitivity variations and B1 field inhomogeneities create intensity variations across the image, making threshold-based methods unreliable.
- **Multiple Contrast Mechanisms:** T1, T2, FLAIR, and diffusion-weighted sequences provide complementary tissue information but require multi-modal fusion strategies.
- **Motion Artifacts:** Patient motion, cardiac pulsation, and respiratory movement introduce blurring and ghosting artifacts.
- **Partial Volume Effects:** Voxels at tissue boundaries contain mixed tissue signals, creating uncertainty in boundary delineation.

Computed Tomography (CT):

- **High Dynamic Range:** CT images span wide intensity ranges (air: -1000 HU, bone: +3000 HU), requiring robust normalization strategies.
- **Noise and Streaking Artifacts:** High noise levels and beam hardening artifacts complicate segmentation.
- **Metal Artifacts:** Dental fillings, surgical implants, and orthopedic hardware create severe streaking artifacts.

2.2.2 Anatomical and Biological Challenges

- **Anatomical Variability:** Natural anatomical variations between individuals complicate atlas-based and template-matching approaches.
- **Pathological Variations:** Diseases alter tissue appearance, texture, and morphology, requiring segmentation methods robust to pathological changes.
- **Tissue Property Variations:** Age, gender, and physiological state influence tissue contrast and appearance.
- **Complex Topologies:** Many anatomical structures have complex, non-convex shapes with thin protrusions and concavities.

2.2.3 Technical and Computational Challenges

- **Class Imbalance:** Background pixels often constitute 80-95% of medical images, biasing models toward majority class prediction.
- **Spatial Resolution Constraints:** Medical images are often acquired at lower resolutions than natural images, limiting fine detail capture.
- **3D Context Requirements:** Many segmentation tasks require understanding 3D spatial relationships and connectivity.
- **Real-time Requirements:** Some applications (e.g., interventional radiology) require real-time segmentation performance.

2.3 Brain MRI Segmentation: A Representative Challenge

Among various medical segmentation tasks, brain MRI segmentation exemplifies many of the challenges described above and serves as an excellent testbed for segmentation algorithms. Brain tissue segmentation involves:

2.3.1 Specific Challenges in Brain Segmentation

- **Complex Anatomy:** The brain contains over 100 distinct anatomical regions with highly variable shapes and sizes.
- **Tissue Contrast Variations:** White matter, gray matter, and cerebrospinal fluid have subtle intensity differences that vary with acquisition parameters and pathology.
- **Intensity Inhomogeneities:** B1 field variations create intensity gradients across brain volumes.
- **Partial Volume Effects:** Thin cortical regions and small subcortical structures suffer from partial volume averaging.
- **Motion Sensitivity:** CSF pulsation and involuntary patient motion create artifacts.

2.3.2 Clinical Relevance

Brain segmentation enables:

- Quantitative assessment of brain atrophy in neurodegenerative diseases
- Surgical planning for tumor resection and epilepsy surgery
- Monitoring of disease progression and treatment response
- Research into brain development, aging, and disease mechanisms

2.4 Research Objectives and Contributions

This assignment focuses on developing and evaluating deep learning solutions for brain tissue segmentation using the IBSR (Internet Brain Segmentation Repository) dataset. Our primary objectives encompass both methodological development and comprehensive evaluation:

2.4.1 Core Objectives

1. **Architecture Implementation:** Develop a robust U-Net architecture optimized for medical image segmentation, including proper initialization, regularization, and training strategies.
2. **Attention Mechanism Integration:** Extend the baseline U-Net with attention gates to improve focus on relevant anatomical structures while suppressing background noise.
3. **Data Processing Pipeline:** Create a comprehensive preprocessing pipeline that addresses data augmentation, normalization, and patch extraction for efficient training.
4. **Loss Function Design:** Implement and evaluate multiple loss functions (Dice, Cross-Entropy, combined) to handle class imbalance and boundary precision requirements.
5. **Rigorous Evaluation:** Conduct thorough quantitative evaluation using multiple metrics and qualitative analysis through comprehensive visualization.
6. **Comparative Analysis:** Systematically compare different architectural variants to understand the impact of design choices.

2.4.2 Technical Contributions

This work contributes to the medical image segmentation field through:

- **Practical Implementation:** A complete, reproducible pipeline from raw data to clinical-grade segmentations.
- **Architecture Analysis:** Detailed comparison of U-Net and Attention U-Net performance characteristics.
- **Loss Function Insights:** Empirical evaluation of different loss functions for imbalanced medical segmentation.
- **Preprocessing Strategies:** Systematic evaluation of data augmentation and preprocessing techniques.
- **Performance Benchmarking:** Comprehensive benchmarking against literature baselines on the IBSR dataset.

2.4.3 Broader Impact

The methodologies and insights developed in this work have implications beyond the specific IBSR dataset:

- **Clinical Translation:** The validated pipeline can be adapted for other brain segmentation tasks and anatomical regions.
- **Research Applications:** The systematic evaluation framework can guide future architecture developments.
- **Educational Value:** The detailed implementation serves as a reference for deep learning in medical imaging.
- **Open Science:** Complete code and methodology enable reproducible research and technology transfer.

2.5 Report Structure

This report is organized to provide a comprehensive treatment of the brain segmentation problem and our solution approach:

1. **Related Work:** Historical context and evolution of segmentation methods, with focus on deep learning approaches.
2. **Dataset and Preprocessing:** Detailed description of the IBSR dataset and our preprocessing pipeline.
3. **Methodology:** Technical details of U-Net and Attention U-Net architectures, loss functions, and evaluation metrics.
4. **Experimental Setup:** Implementation details, hyperparameters, and training procedures.
5. **Results and Analysis:** Comprehensive quantitative and qualitative evaluation of model performance.
6. **Discussion:** Interpretation of results, limitations, and future directions.
7. **Conclusion:** Summary of contributions and clinical implications.

The U-Net architecture, originally proposed by Ronneberger et al. [?], has become the de facto standard for medical image segmentation due to its ability to capture both local and global context through its symmetric encoder-decoder structure with skip connections. The Attention U-Net, introduced by Oktay et al. [?], enhances this architecture by incorporating attention gates that learn to focus on relevant anatomical structures while suppressing irrelevant background information. Our implementation and evaluation of these architectures forms the core technical contribution of this work.

3 Related Work

The field of medical image segmentation has undergone a dramatic transformation from traditional computer vision techniques to sophisticated deep learning approaches. This evolution reflects both technological advances and the growing recognition of segmentation as a fundamental component of computer-aided diagnosis and treatment planning. This section provides a comprehensive analysis of this evolution, with particular emphasis on the technical innovations that have shaped the current state-of-the-art.

3.1 Traditional Segmentation Methods: Foundations and Limitations

Before the deep learning revolution, medical image segmentation relied on classical computer vision and image processing techniques. While these methods laid crucial groundwork, their limitations highlighted the need for more sophisticated approaches capable of handling the complexity of medical imaging.

3.1.1 Thresholding-Based Methods

Thresholding represents the simplest class of segmentation algorithms, operating on the principle that different tissues can be distinguished by their intensity values.

Global Thresholding:

- **Otsu’s Method** [?]: Automatically determines optimal threshold by maximizing inter-class variance. Computationally efficient but assumes bimodal intensity distributions.
- **Multi-level Thresholding**: Extends Otsu to multiple classes using iterative optimization or evolutionary algorithms.

Adaptive Thresholding:

- **Local Thresholding**: Computes thresholds for local image regions, addressing intensity inhomogeneities.
- **Bernsen’s Method** [?]: Uses local contrast to determine thresholds, effective for images with varying illumination.

Limitations in Medical Imaging:

- **Intensity Inhomogeneities**: MRI images often exhibit smooth intensity variations due to coil sensitivity profiles and B1 field inhomogeneities.
- **Multi-modal Distributions**: Medical images frequently have complex intensity distributions that don’t conform to simple thresholding assumptions.
- **Noise Sensitivity**: Thresholding methods are highly sensitive to noise and acquisition artifacts.
- **Shape Constraints**: No inherent mechanisms to enforce anatomical shape constraints or spatial relationships.

3.1.2 Region-Based Methods

Region-based approaches grow or merge image regions based on similarity criteria, offering more sophisticated segmentation than simple thresholding.

Region Growing:

- **Seed-Based Growth:** Starts from user-selected seed points and expands to neighboring pixels within intensity similarity thresholds.
- **Connected Component Analysis:** Groups connected pixels with similar properties.
- **Hybrid Approaches:** Combine region growing with edge detection for improved boundary accuracy.

Region Splitting and Merging:

- **Quadtree Decomposition:** Recursively splits image into quadrants until homogeneity criteria are met.
- **Merge Criteria:** Adjacent regions are merged based on statistical similarity measures (e.g., mean intensity differences).

Mathematical Formulation: The region growing process can be formulated as:

$$R_{n+1} = R_n \cup \{p \in \mathcal{N}(R_n) \mid |I(p) - \mu(R_n)| < T\} \quad (1)$$

where R_n is the region at step n , $\mathcal{N}(R_n)$ are neighboring pixels, $I(p)$ is pixel intensity, $\mu(R_n)$ is region mean, and T is the similarity threshold.

Clinical Applications and Limitations:

- **Advantages:** Can handle complex shapes and incorporate texture information.
- **Seed Dependency:** Requires manual seed selection, limiting automation potential.
- **Leakage Issues:** May "leak" into adjacent tissues with similar intensities.
- **Parameter Sensitivity:** Performance heavily depends on threshold and similarity parameters.

3.1.3 Deformable Models and Active Contours

Active contour models, commonly known as "snakes," represent boundaries as parametric curves that evolve to minimize energy functionals.

Classical Snake Model [?]: The snake energy functional combines internal and external energies:

$$E_{\text{snake}} = \int_0^1 [E_{\text{int}}(v(s)) + E_{\text{ext}}(v(s))] ds \quad (2)$$

where $v(s) = [x(s), y(s)]$ parameterizes the contour, E_{int} enforces smoothness constraints, and E_{ext} attracts the contour to image features.

Geodesic Active Contours [?]: Reformulate active contours using level set methods:

$$\frac{\partial \phi}{\partial t} = |\nabla \phi| \cdot \left(\alpha \cdot \operatorname{div} \left(\frac{\nabla \phi}{|\nabla \phi|} \right) - \beta - \gamma \cdot \kappa \right) \quad (3)$$

where ϕ is the level set function, κ is curvature, and α, β, γ are weighting parameters.

Gradient Vector Flow (GVF) Snakes [?]: Improve convergence for concave boundaries by using gradient vector flow fields:

$$\nabla^2 \mathbf{v} - \mathbf{v} + \nabla f = 0 \quad (4)$$

Advantages and Applications:

- **Boundary Focus:** Explicitly model object boundaries rather than regions.
- **Shape Flexibility:** Can handle complex, non-convex shapes.
- **Clinical Success:** Widely used in cardiac segmentation, vessel tracking, and organ boundary delineation.

Limitations:

- **Initialization Sensitivity:** Require good initial contours close to true boundaries.
- **Topology Constraints:** Struggle with topological changes (splitting, merging).
- **Parameter Tuning:** Multiple parameters (α, β, γ) require careful optimization.
- **Computational Cost:** Iterative optimization can be time-consuming for high-resolution images.

3.1.4 Graph-Based Methods

Graph-based approaches formulate segmentation as graph partitioning problems, offering principled optimization frameworks.

Graph Cuts [?]:

- Model image as graph where pixels are nodes and edges represent pixel similarities.
- Find minimum cut separating foreground and background seeds.
- **Max-Flow/Min-Cut Theorem:** Guarantees optimal solution for binary segmentation.

Random Walker Segmentation [?]:

- Model pixel labels as probabilities from random walks starting at seeded pixels.
- Solve harmonic equation: $\mathbf{L}\mathbf{x} = \mathbf{b}$
- where $\mathbf{L}$ is graph Laplacian, $\mathbf{x}$ are label probabilities, $\mathbf{b}$ are boundary conditions.

Mathematical Framework: The graph cut energy function is:

$$E(f) = \sum_{p \in \mathcal{P}} R_p(f_p) + \sum_{(p,q) \in \mathcal{N}} B_{pq}(f_p, f_q) \quad (5)$$

where R_p is regional term, B_{pq} is boundary term, and f_p assigns labels to pixels.

Advantages:

- **Global Optimization:** Provide theoretical guarantees of optimality.
- **Shape Priors:** Can incorporate anatomical constraints through regional terms.
- **Multi-Object Segmentation:** Naturally extend to multiple classes.

Limitations in Medical Imaging:

- **Seed Requirements:** Need good foreground/background seeds.
- **Scalability Issues:** Computational complexity grows rapidly with image size.
- **Feature Design:** Performance depends on quality of boundary and regional features.

3.2 The Deep Learning Revolution in Medical Segmentation

The advent of deep learning, particularly convolutional neural networks, fundamentally transformed medical image segmentation by automating feature learning and providing end-to-end trainable systems.

3.2.1 Fully Convolutional Networks (FCNs): The Foundation

Long et al. [?] pioneered dense prediction with CNNs by replacing fully connected layers with convolutional layers, enabling pixel-wise predictions.

Architecture Innovations:

- **Encoder-Decoder Structure:** Downsampling encoder captures semantic information, upsampling decoder recovers spatial resolution.
- **Skip Connections:** Fuse features from different scales to combine semantic and spatial information.
- **Transfer Learning:** Fine-tune networks pre-trained on ImageNet for medical tasks.

Technical Details: FCNs use bilinear upsampling and feature concatenation:

$$\mathbf{F}^{(l)} = \mathbf{F}_{\text{upsampled}}^{(l)} \oplus \mathbf{F}_{\text{encoder}}^{(l)} \quad (6)$$

where $\oplus$ denotes concatenation and l indicates resolution level.

Medical Applications: FCNs achieved state-of-the-art results on multiple datasets but suffered from coarse predictions due to upsampling limitations.

3.2.2 U-Net: A Breakthrough in Medical Imaging

The U-Net architecture [?] revolutionized medical image segmentation through its elegant encoder-decoder design specifically tailored for biomedical applications.

Architectural Innovations:

- **Symmetric Design:** Balanced encoder-decoder with identical depth.
- **Skip Connections:** Direct concatenation of encoder features to decoder at each resolution level.
- **Feature Preservation:** Maintains high-resolution features throughout the network.
- **Data Efficiency:** Achieves excellent performance with limited training data.

Mathematical Formulation: The U-Net forward pass combines convolution, down-sampling, and upsampling operations:

$$\mathbf{x}_{\text{encoder}}^{(l+1)} = \text{MaxPool}(\text{ReLU}(\text{Conv}(\mathbf{x}_{\text{encoder}}^{(l)}))) \quad (7)$$

$$\mathbf{x}_{\text{decoder}}^{(l)} = \text{ReLU}(\text{Conv}(\text{Up}(\mathbf{x}_{\text{decoder}}^{(l+1)}) \oplus \mathbf{x}_{\text{encoder}}^{(l)})) \quad (8)$$

Impact and Success:

- **Widespread Adoption:** Became the standard architecture for medical segmentation.
- **Challenge Dominance:** Won multiple medical imaging challenges (ISBI cell tracking, Kaggle ultrasound nerve segmentation).
- **Versatility:** Successfully applied to diverse organs (liver, prostate, brain, retina).
- **Extensions:** Inspired numerous variants (V-Net, U-Net++, Attention U-Net).

3.2.3 3D Extensions and Volumetric Segmentation

Direct 3D segmentation addresses the limitations of 2D approaches by modeling spatial relationships in three dimensions.

V-Net Architecture [?]:

- **3D Operations:** Uses 3D convolutions, max pooling, and transposed convolutions.
- **Dice Loss:** Introduces Dice coefficient as loss function for class imbalance.
- **Residual Connections:** Incorporates residual blocks for improved gradient flow.

3D U-Net Variants:

- **No New-Net [?]:** Optimized 3D U-Net with instance normalization and leaky ReLU.
- **Progressive U-Net:** Multi-stage refinement for improved accuracy.
- **Hybrid Approaches:** Combine 2D and 3D operations for computational efficiency.

Challenges and Solutions:

- **Memory Constraints:** 3D operations require significant GPU memory.
- **Data Requirements:** Need large annotated 3D datasets.
- **Computational Cost:** Training time increases dramatically with volume size.

3.2.4 Atrous Convolutions and Multi-Scale Processing

Dilated convolutions enable large receptive fields without losing resolution, crucial for capturing multi-scale anatomical structures.

DeepLab Architecture [?]:

- **Atrous Convolutions:** Dilated filters capture multi-scale context.
- **Atrous Spatial Pyramid Pooling (ASPP):** Parallel dilated convolutions at multiple rates.
- **Conditional Random Fields (CRF):** Refine predictions with spatial consistency.

Mathematical Foundation: A dilated convolution is defined as:

$$y[i] = \sum_k x[i + r \cdot k] \cdot w[k] \quad (9)$$

where r is the dilation rate, enabling exponential receptive field growth.

Medical Applications:

- **Retinal Vessel Segmentation:** Captures vessels at multiple scales.
- **Lung Nodule Detection:** Handles nodules of varying sizes.
- **Brain Lesion Segmentation:** Processes lesions with complex morphologies.

3.3 Attention Mechanisms: Focus and Selectivity

Attention mechanisms, originally developed for sequence processing, have been successfully adapted to computer vision tasks, particularly medical image segmentation.

3.3.1 Attention U-Net: Learning Where to Look

Oktay et al. [?] introduced attention gates to U-Net, enabling selective feature aggregation.

Attention Gate Mechanism:

- **Gating Signal:** Features from decoder provide contextual information.
- **Attention Coefficients:** Learned weights determine feature relevance.
- **Soft Attention:** Sigmoid-activated coefficients enable soft selection.

Mathematical Implementation: The attention coefficient computation:

$$\alpha_i = \sigma \left(W_\psi^T \cdot \tanh \left(W_g^T g + W_x^T x_i \right) \right) \quad (10)$$

where g is gating signal, x_i are input features, and W are learned weight matrices.

Weighted Feature Aggregation:

$$\mathbf{x}'_i = \alpha_i \cdot x_i \quad (11)$$

Benefits in Medical Imaging:

- **Background Suppression:** Reduces false positives in background regions.
- **Boundary Enhancement:** Improves accuracy at tissue boundaries.
- **Robustness:** Better handling of anatomical variations.

3.3.2 Self-Attention and Transformer Architectures

Recent work explores self-attention mechanisms for global context modeling in medical segmentation.

Vision Transformers (ViT):

- **Patch Embeddings:** Divide images into patches and embed in high-dimensional space. **Multi-Head Self-Attention:** Model long-range dependencies across image regions.
- **Hierarchical Processing:** Combine local and global feature learning.

Medical Applications:

- **Swin Transformer:** Hierarchical attention for multi-scale feature extraction.
- **Medical ViT Variants:** Adapted for limited medical imaging data.
- **Hybrid Approaches:** Combine CNNs with transformers for optimal performance.

3.3.3 Spatial and Channel Attention

Various attention mechanisms focus on different aspects of feature modulation.

Squeeze-and-Excitation (SE) Blocks [?]:

- **Channel Attention:** Learn per-channel importance weights.
- **Global Average Pooling:** Aggregate spatial information.
- **Feature Recalibration:** Rescale channels based on global context.

CBAM: Convolutional Block Attention Module [?]:

- **Spatial Attention:** Focus on important spatial locations.
- **Channel Attention:** Emphasize informative channels.
- **Multiplicative Integration:** Combine spatial and channel attention.

3.4 Current State-of-the-Art and Emerging Trends

3.4.1 Multi-Scale and Multi-Resolution Architectures

Modern architectures increasingly incorporate multi-scale processing to handle anatomical structures of varying sizes.

HRNet: High-Resolution Networks [?]:

- **Parallel Multi-Resolution:** Maintain high-resolution representations throughout.
- **Resolution Fusion:** Exchange information across different scales.
- **Memory Efficiency:** Avoid repeated downsampling and upsampling.

Feature Pyramid Networks (FPN):

- **Top-Down Pathway:** Propagate semantic information to higher resolution.
- **Lateral Connections:** Fuse features across scales.
- **Unified Representation:** Produce scale-aware feature maps.

3.4.2 Ensemble and Uncertainty Estimation

Combining multiple models improves robustness and enables uncertainty quantification.

Ensemble Methods:

- **Model Averaging:** Combine predictions from diverse architectures.
- **Bootstrap Aggregation:** Train on different data subsets.
- **Test-Time Augmentation:** Average predictions across augmentations.

Uncertainty Quantification:

- **Monte Carlo Dropout:** Estimate uncertainty through multiple forward passes.
- **Bayesian Neural Networks:** Model weight uncertainty.
- **Deep Ensembles:** Multiple models provide uncertainty estimates.

3.4.3 Transfer Learning and Domain Adaptation

Leveraging pre-trained models addresses limited medical data availability.

Domain-Specific Pre-training:

- **Medical ImageNet:** Pre-train on diverse medical imaging datasets.
- **Self-Supervised Learning:** Learn representations without manual annotations.
- **Contrastive Learning:** Pull similar images together, push dissimilar apart.

Domain Adaptation Techniques:

- **Adversarial Training:** Align feature distributions across domains.
- **Cycle-Consistent Learning:** Learn domain-invariant representations.
- **Style Transfer:** Normalize appearance variations.

3.4.4 Computational Efficiency and Deployment

Recent trends focus on making segmentation models practical for clinical deployment.

Model Compression:

- **Knowledge Distillation:** Train compact models to mimic larger networks.
- **Pruning:** Remove redundant parameters.
- **Quantization:** Reduce numerical precision.

Efficient Architectures:

- **MobileNets:** Depthwise separable convolutions for mobile deployment.
- **ShuffleNets:** Channel shuffling for efficient feature mixing.
- **EfficientNets:** Compound scaling for optimal accuracy-efficiency trade-offs.

3.5 Summary and Positioning of Current Work

The evolution from traditional methods to modern deep learning approaches represents a paradigm shift in medical image segmentation. Early methods relied on handcrafted features and optimization-based approaches, while contemporary methods leverage data-driven learning for automatic feature discovery.

Our work builds upon this foundation by implementing and evaluating U-Net and Attention U-Net architectures for brain tissue segmentation. These architectures represent the current practical state-of-the-art, balancing performance with computational feasibility. The systematic evaluation provides insights into architectural trade-offs and establishes baselines for future improvements.

Key positioning of our contribution:

- **Architectural Focus:** Comprehensive analysis of U-Net variants rather than novel architecture design.
- **Practical Implementation:** Production-ready code following best practices.
- **Thorough Evaluation:** Multi-metric assessment with detailed error analysis.
- **Reproducibility:** Complete methodology enabling replication and extension.

The subsequent sections detail our implementation, evaluation, and findings within this rich context of segmentation research.

4 Dataset and Preprocessing

4.1 Dataset Description

The IBSR (Internet Brain Segmentation Repository) dataset represents one of the most extensively studied and validated resources for evaluating brain tissue segmentation algorithms. This dataset has served as a benchmark for numerous segmentation methods and continues to be a cornerstone of medical image segmentation research.

4.1.1 Dataset Overview and Acquisition

The IBSR dataset was created through a collaborative effort between Massachusetts General Hospital and the Surgical Planning Laboratory at Brigham and Women’s Hospital. The dataset comprises high-resolution T1-weighted MRI scans acquired using standardized protocols to ensure consistency across subjects.

Acquisition Parameters:

- **Scanner:** General Electric 1.5 Tesla Signa scanner
- **Sequence:** 3D spoiled gradient echo (SPGR) T1-weighted
- **TR/TE:** 35 ms / 5 ms (repetition time / echo time)
- **Flip Angle:** 45 degrees
- **Slice Thickness:** 1.5 mm (overlapping slices)
- **In-plane Resolution:** 256 × 256 pixels
- **Field of View:** 240 × 240 mm
- **Voxel Dimensions:** 0.9375 × 0.9375 × 1.5 mm³

4.1.2 Subject Demographics and Inclusion Criteria

The dataset includes MRI scans from 18 neurologically normal subjects, providing a representative sample of adult brain anatomy:

Demographic Characteristics:

- **Age Range:** 21-50 years (mean age: 35.2 years)
- **Gender Distribution:** 9 male, 9 female subjects
- **Health Status:** Neurologically normal adults with no history of neurological disorders
- **Inclusion Criteria:** No visible anatomical abnormalities or artifacts that would compromise segmentation quality

This demographic composition ensures anatomical variability while maintaining the dataset’s focus on normal brain anatomy rather than pathological cases.

4.1.3 Ground Truth Segmentation Protocol

The ground truth segmentations were created through a rigorous manual annotation process that represents the gold standard for brain tissue segmentation:

Annotation Methodology:

- **Expert Annotators:** Segmentations performed by trained neuroanatomists with expertise in brain anatomy
- **Quality Control:** Multiple rounds of review and consensus resolution for disputed regions

- **Inter-rater Reliability:** Assessment of consistency between different annotators
- **Resolution:** Full 3D voxel-level annotations rather than slice-by-slice segmentation

Segmentation Classes: The dataset provides segmentations for major brain tissue classes:

1. **Cerebrospinal Fluid (CSF):** Ventricular spaces and subarachnoid spaces
2. **Gray Matter (GM):** Cerebral cortex and deep gray matter structures
- 3.
4. **White Matter (WM):** Myelinated axons and fiber tracts
5. **Mixed/Partial Volume:** Voxels containing multiple tissue types

For binary segmentation tasks (as in this assignment), the mixed/partial volume class is typically merged with background or handled through specialized preprocessing.

4.1.4 Data Format and Structure

The dataset is distributed in NIfTI format, the standard for medical imaging data:

File Organization:

- **Volume Files:** T1-weighted anatomical images (*.nii files)
- **Segmentation Files:** Corresponding tissue segmentations (*.nii files)
- **Naming Convention:** IBSR_XX format (XX = subject number from 01 to 18)
- **Co-registration:** Volume and segmentation files are perfectly aligned

Volume Characteristics:

- **Dimensions:** Typically 256 (E) 256 (E) 128 voxels (varies slightly by subject)
- **Data Type:** 16-bit integer values representing MRI signal intensity
- **Orientation:** Standard radiological orientation (RAS coordinate system)
- **Intensity Range:** Raw values range from 0 to approximately 4000 (scanner-dependent)

4.1.5 Technical Challenges and Dataset Characteristics

The IBSR dataset exhibits several characteristics that make it both valuable for research and challenging for segmentation algorithms:

Intensity Characteristics:

- **Dynamic Range:** Wide intensity variations across tissues (CSF: low, WM: high, GM: intermediate)
- **Contrast Properties:** Excellent gray-white matter contrast due to T1-weighting
- **Inhomogeneity:** Gradual intensity variations due to coil sensitivity profiles

- **Noise Level:** Minimal acquisition noise due to high SNR protocol

Anatomical Variability:

- **Size Variations:** Brain volumes range from approximately 1100-1600 cm³
- **Shape Variations:** Natural anatomical differences in gyral patterns and sulcal depths
- **Ventricular Size:** Variations in ventricular volume across subjects
- **Cortical Thickness:** Individual variations in cortical mantle thickness

Technical Artifacts:

- **Motion Artifacts:** Minimal due to cooperative adult subjects
- **Flow Artifacts:** Potential artifacts in ventricular regions due to CSF flow
- **Partial Volume Effects:** Significant at tissue boundaries, especially cortex/CSF interfaces
- **Gibbs Ringing:** Possible truncation artifacts at high-contrast boundaries

4.1.6 Data Processing Strategy

Our approach transforms the 3D volumetric data into a format suitable for 2D deep learning models while maximizing information utilization:

3D to 2D Conversion:

- **Multi-Planar Reformation:** Extract slices along three orthogonal planes (axial, coronal, sagittal)
- **Isotropic Sampling:** Ensure consistent sampling across all orientations
- **Complete Coverage:** Process all slices to capture full anatomical information

Patch-Based Processing:

- **Patch Size:** 128x128 pixels balances receptive field requirements with memory constraints
- **Patch Extraction:** Four non-overlapping patches per slice maximize data utilization
- **Spatial Context:** Patches maintain local anatomical relationships

Data Augmentation Through Multi-Axial Processing:

- **Orientation Diversity:** Three orthogonal views provide anatomical context from different perspectives
- **Scale Invariance:** Different orientations capture structures at various apparent sizes
- **Artifact Robustness:** Multiple views reduce sensitivity to orientation-specific artifacts

4.1.7 Dataset Statistics and Analysis

Comprehensive statistical analysis reveals important characteristics for algorithm development:

Volume Statistics:

- **Total Voxels:** Approximately 8.4 million voxels per volume
- **Tissue Distribution:** CSF (15-20%), GM (40-45%), WM (30-35%)
- **Surface Area:** Cortical surface area ranges from 1500-2000 cm²
- **Fold Complexity:** Mean cortical folding index varies by individual anatomy

Challenge Assessment:

- **Boundary Complexity:** Extensive gray-white matter interfaces with complex folding patterns
- **Thin Structures:** Subcortical nuclei and cortical regions pose segmentation challenges
- **Connectivity:** White matter tracts require 3D connectivity preservation
- **Symmetry:** Bilateral brain structures provide natural consistency checks

4.1.8 Comparison with Other Brain Segmentation Datasets

The IBSR dataset occupies a unique position among brain segmentation benchmarks:

Advantages over Other Datasets:

- **Expert Annotations:** High-quality manual segmentations by trained anatomists
- **Complete Coverage:** Full brain volumes rather than limited regions
- **Standardization:** Consistent acquisition protocol across subjects
- **Size:** Sufficient subjects for meaningful statistical analysis

Limitations:

- **Scale:** Smaller than modern datasets (e.g., ADNI, UK Biobank)
- **Demographics:** Limited age range and demographic diversity
- **Pathology:** Only normal anatomy, no disease cases
- **Resolution:** Lower resolution compared to ultra-high-field imaging

Positioning in Research Landscape: The IBSR dataset serves as an excellent benchmark for methodological development and algorithm validation, particularly for:

- **Methodological Research:** Testing novel segmentation algorithms
- **Comparative Studies:** Establishing performance baselines
- **Technical Development:** Optimizing preprocessing and postprocessing techniques
- **Educational Purposes:** Learning segmentation concepts with high-quality data

4.1.9 Dataset Utilization in This Work

Our implementation leverages the IBSR dataset’s characteristics for robust model development:

Training Strategy:

- **Cross-Validation:** Volume-level splitting ensures unbiased evaluation
- **Multi-Planar Training:** Utilize all three anatomical planes for comprehensive learning
- **Patch Diversity:** Extract multiple patches per slice to increase sample diversity

Evaluation Considerations:

- **Inter-Subject Variability:** Account for natural anatomical differences
- **Annotation Quality:** Leverage high-quality ground truth for reliable assessment
- **Clinical Relevance:** Focus on metrics relevant to clinical applications

This comprehensive dataset forms the foundation for our segmentation experiments, providing both the challenges and the validation necessary for developing robust medical image segmentation algorithms.

4.2 Data Preprocessing Pipeline

The preprocessing pipeline transforms raw 3D MRI volumes into standardized 2D patches suitable for deep learning training. This multi-stage process addresses data format conversion, augmentation, normalization, and memory optimization while preserving anatomical information and spatial relationships.

4.2.1 3D to 2D Slice Extraction Strategy

The conversion from 3D volumes to 2D slices represents a critical design decision balancing computational efficiency with anatomical information preservation.

Multi-Planar Reformation Rationale:

- **Anatomical Diversity:** Three orthogonal planes (axial, coronal, sagittal) capture brain structures from complementary viewpoints
- **Scale Invariance:** Structures appear at different sizes across orientations, improving model robustness
- **Artifact Distribution:** Orientation diversity reduces sensitivity to plane-specific artifacts
- **Data Augmentation:** Natural augmentation through multi-planar viewing without artificial transformations

Slice Selection Parameters:

- **Starting Index (10):** Avoids edge artifacts and skull-stripping boundary effects common in the first slices

- **Stride (3)**: Provides adequate slice diversity while maintaining computational tractability
- **Maximum Slices (48)**: Balances dataset size with anatomical coverage (approximately 72mm coverage at 1.5mm slice thickness)

Mathematical Formulation: For a volume of dimensions (D_x, D_y, D_z) , slices are extracted as:

$$\text{Axial slices : } \{V[:, :, k] \mid k = 10, 13, 16, \dots, \min(D_z - 10, 154)\} \quad (12)$$

$$\text{Coronal slices : } \{V[:, k, :] \mid k = 10, 13, 16, \dots, \min(D_y - 10, 154)\} \quad (13)$$

$$\text{Sagittal slices : } \{V[k, :, :] \mid k = 10, 13, 16, \dots, \min(D_x - 10, 154)\} \quad (14)$$

4.2.2 Orientation Standardization and Rotation

Standardizing slice orientation ensures consistent anatomical positioning across all training samples.

Rotation Operation:

- **90-degree Counterclockwise:** Aligns brain anatomy with conventional radiological viewing
- **Consistent Orientation:** Ensures left-right symmetry and anatomical landmarks appear in expected positions
- **Network Consistency:** Prevents confusion from orientation variations during training

Mathematical Implementation: The rotation operation transforms coordinates (x, y) to (x', y') :

$$\begin{pmatrix} x' \\ y' \end{pmatrix} = \begin{pmatrix} \cos \theta & -\sin \theta \\ \sin \theta & \cos \theta \end{pmatrix} \begin{pmatrix} x \\ y \end{pmatrix}, \quad \theta = \pi/2 \quad (15)$$

4.2.3 Padding Strategy for Spatial Consistency

Padding ensures uniform input dimensions while preserving spatial relationships and minimizing boundary artifacts.

Padding Design Considerations:

- **Target Dimensions (256x256):** Accommodates network input requirements while maintaining reasonable memory usage
- **Zero Padding:** Maintains background intensity consistency and avoids introducing artificial structures
- **Centered Padding:** Preserves relative anatomical positions and spatial relationships

Padding Calculation: For a slice of dimensions (h, w) and target (H, W) :

$$\text{Padding width left/right : } p_w = \lfloor (W - w)/2 \rfloor, \quad p'_w = W - w - p_w \quad (16)$$

$$\text{Padding height top/bottom : } p_h = \lfloor (H - h)/2 \rfloor, \quad p'_h = H - h - p_h \quad (17)$$

Rationale for Zero Padding:

- **Background Consistency:** Zero values match typical background intensities in MRI
- **Boundary Effects:** Minimizes edge artifacts that could affect segmentation near image borders
- **Network Robustness:** Trains network to handle variable field-of-view scenarios

4.2.4 Patch Extraction for Memory and Context Optimization

Patch-based processing enables training on high-resolution images while managing GPU memory constraints.

Patch Design Parameters:

- **Patch Size (128×128):** Balances receptive field requirements with memory constraints
- **Non-overlapping Extraction:** Maximizes unique anatomical information per slice
- **Quadrant Division:** Systematic spatial coverage of the entire slice

Patch Coordinates: The four patches divide the 256×256 padded slice into quadrants:

$$P_1(\text{top-left}) : [0 : 128, 0 : 128] \quad (18)$$

$$P_2(\text{top-right}) : [0 : 128, 128 : 256] \quad (19)$$

$$P_3(\text{bottom-left}) : [128 : 256, 0 : 128] \quad (20)$$

$$P_4(\text{bottom-right}) : [128 : 256, 128 : 256] \quad (21)$$

Technical Advantages:

- **Memory Efficiency:** Smaller patches fit better in GPU memory during training
- **Context Preservation:** 128×128 provides sufficient anatomical context for tissue classification
- **Data Augmentation:** Multiple patches per slice increase dataset diversity
- **Resolution Maintenance:** Preserves native image resolution without downsampling

4.2.5 Segmentation Mask Processing

Specialized processing ensures segmentation masks are compatible with loss functions and network outputs.

Binary Classification Setup:

- **Background (0):** Non-brain tissues and CSF regions
- **Foreground (1):** Brain parenchyma (gray + white matter combined)
- **Data Type:** 64-bit integers for PyTorch compatibility

Mask Binarization: Critical preprocessing step applied in the data loading pipeline:

$$M_{\text{binary}} = (M_{\text{original}} > 0) \in \{0, 1\} \quad (22)$$

Implementation Location: Applied in `IBSRPatchDataset.__getitem__()` to ensure consistency across training and evaluation.

4.2.6 Intensity Normalization Strategy

While not implemented in the current pipeline, intensity normalization represents an important consideration:

Potential Normalization Approaches:

- **Z-score Normalization:** $(I - \mu)/\sigma$ using per-volume statistics
- **Min-Max Scaling:** $(I - I_{\min})/(I_{\max} - I_{\min})$ for fixed range $[0,1]$
- **Robust Scaling:** Using median and IQR for outlier resistance
- **Histogram Matching:** Normalization to reference intensity distribution

Current Approach Rationale: The implementation uses raw intensity values, justified by:

- **Data Integrity:** Preserves original intensity relationships for tissue discrimination
- **Model Capacity:** Allows network to learn intensity mappings from data
- **Consistency:** Matches common practice in medical image segmentation literature

4.2.7 Data Loading and Augmentation Pipeline

The `IBSRPatchDataset` class implements efficient data loading with preprocessing integration:

Constructor Operations:

1. Load all volume and segmentation file pairs
2. Extract slices from three orthogonal planes
3. Generate patch coordinates for each slice
4. Store patch references for on-demand loading

getitem Implementation:

1. Load volume and segmentation slices on-demand
2. Apply rotation and padding operations
3. Extract specific patch based on index
4. Apply mask binarization
5. Convert to PyTorch tensors

Memory Optimization:

- **Lazy Loading:** Patches loaded only when requested, reducing memory footprint
- **Just-in-Time Processing:** Preprocessing applied per batch, enabling dynamic augmentation
- **Cache Efficiency:** Repeated access to same volumes leverages OS caching

4.2.8 Dataset Splitting Strategy

Proper data partitioning ensures unbiased model evaluation and prevents information leakage.

Volume-Level Splitting:

- **Unit of Splitting:** Complete brain volumes rather than individual patches
- **Stratification:** Maintain anatomical diversity across splits
- **Proportions:** 70% training, 15% validation, 15% test

Splitting Algorithm:

1. Shuffle volume indices with fixed random seed
2. Calculate split sizes: $n_{\text{train}} = 0.7N$, $n_{\text{val}} = 0.15N$, $n_{\text{test}} = 0.15N$
3. Assign volumes to splits sequentially

Leakage Prevention:

- **Volume Separation:** No patches from training volumes appear in validation/test
- **Spatial Separation:** Different anatomical regions across splits
- **Subject Independence:** Complete subject separation prevents subject-specific bias

4.2.9 Preprocessing Quality Control

Several mechanisms ensure preprocessing quality and reproducibility:

Validation Checks:

- **Dimension Verification:** Confirm all patches have correct 128E128 dimensions
- **Mask Consistency:** Verify segmentation masks contain only binary values
- **Data Type Validation:** Ensure proper tensor types for PyTorch compatibility
- **Range Checking:** Validate intensity values are within expected ranges

Reproducibility Measures:

- **Fixed Random Seeds:** Deterministic slice and patch selection
- **Order Preservation:** Consistent volume processing order
- **Parameter Documentation:** All preprocessing parameters explicitly documented

4.2.10 Computational Performance Analysis

The preprocessing strategy achieves optimal trade-offs between data utilization and computational efficiency:

Data Augmentation Factor:

- **Slice Factor:** 48 slices $\times$ 3 axes = 144 slices per volume
- **Patch Factor:** 4 patches per slice = 576 patches per volume
- **Total Augmentation:** 576 $\times$ increase in training samples from single volume

Memory Efficiency:

- **Per-Patch Memory:** 128 $\times$ 128 $\times$ 4 bytes = 65KB per patch
- **Batch Memory:** 16 patches $\times$ 65KB = 1MB per batch
- **GPU Utilization:** Efficient memory usage enables larger batch sizes

Training Time Considerations:

- **I/O Overhead:** File loading represents primary bottleneck
- **CPU Utilization:** Preprocessing uses CPU cores efficiently
- **GPU Utilization:** Minimal preprocessing on GPU maximizes training efficiency

This comprehensive preprocessing pipeline transforms raw medical imaging data into a format optimized for deep learning while preserving anatomical information and ensuring computational efficiency.

4.2.11 Dataset Split

The dataset is split into:

- Training set: 70% of volumes
- Validation set: 15% of volumes
- Test set: 15% of volumes

This split ensures sufficient training data while maintaining separate validation and test sets for unbiased evaluation.

4.3 Dataset Statistics

Comprehensive dataset analysis is crucial for understanding data characteristics and potential challenges. Figure ?? shows the dataset statistics including:

- **Image intensity distribution:** Understanding the pixel value distribution helps in normalization and preprocessing decisions. The histogram reveals the range and distribution of intensity values, which is important for proper normalization.